

# ADS\_505\_Final\_Project

October 9, 2025

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, \
    HistGradientBoostingClassifier
from sklearn.svm import LinearSVC
from sklearn.neural_network import MLPClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from sklearn.metrics import classification_report, confusion_matrix, \
    roc_auc_score, roc_curve, average_precision_score, precision_recall_curve
from scipy.stats import randint, uniform

%matplotlib inline
```

```
[2]: df = pd.read_csv('creditcard.csv')
df.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7 \
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941

	V8	V9	...	V21	V22	V23	V24	V25 \
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170

```

2  0.247676 -1.514654 ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024 ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739 ... -0.009431  0.798278 -0.137458  0.141267 -0.206010

```

```

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053  149.62      0
1  0.125895 -0.008983  0.014724   2.69      0
2 -0.139097 -0.055353 -0.059752  378.66      0
3 -0.221929  0.062723  0.061458  123.50      0
4  0.502292  0.219422  0.215153   69.99      0

```

[5 rows x 31 columns]

Checking data frame. Data only contains numerical inputs from V1 to V28

## 1 EDA

```

[3]: # Distribution of fraud class
sns.countplot(x = 'Class', data = df);
plt.title('Distribution of Classes (0 = No Fraud ; 1 = Fraud)');
plt.show


print(df['Class'].value_counts(normalize = True)) # percentage of fraud and no_
↪ fraud cases

```

```

Class
0    0.998273
1    0.001727
Name: proportion, dtype: float64

```



Fraud cases are extremely rare (only 0.172% of all transactions)

```
[4]: print(df.describe())
      print(df.shape)
```

	Time	V1	V2	V3	V4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	V5	V6	V7	V8	V9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16	-2.406331e-15
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01

50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	...	V21	V22	V23	V24 \
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	1.654067e-16	-3.568593e-16	2.578648e-16	4.473266e-15
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	5.340915e-16	1.683437e-15	-3.660091e-16	-1.227390e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 31 columns]  
(284807, 31)

```
[5]: df.info()
      print(df.duplicated().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Time    284807 non-null    float64
1    V1       284807 non-null    float64
2    V2       284807 non-null    float64
```

```

3   V3      284807 non-null float64
4   V4      284807 non-null float64
5   V5      284807 non-null float64
6   V6      284807 non-null float64
7   V7      284807 non-null float64
8   V8      284807 non-null float64
9   V9      284807 non-null float64
10  V10     284807 non-null float64
11  V11     284807 non-null float64
12  V12     284807 non-null float64
13  V13     284807 non-null float64
14  V14     284807 non-null float64
15  V15     284807 non-null float64
16  V16     284807 non-null float64
17  V17     284807 non-null float64
18  V18     284807 non-null float64
19  V19     284807 non-null float64
20  V20     284807 non-null float64
21  V21     284807 non-null float64
22  V22     284807 non-null float64
23  V23     284807 non-null float64
24  V24     284807 non-null float64
25  V25     284807 non-null float64
26  V26     284807 non-null float64
27  V27     284807 non-null float64
28  V28     284807 non-null float64
29  Amount  284807 non-null float64
30  Class   284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
1081

```

No missing values or duplicate rows found.

## 2 Correlation Heat Map

```

[6]: # Correlation heat map to see if any variables are highly correlated to class
correlation_mat = df.corr()

plt.figure(figsize = (30,20))
sns.heatmap(correlation_mat,cmap = 'coolwarm', annot = True, center = 0)

corr_with_class = correlation_mat['Class'].sort_values(ascending = False)
print(corr_with_class)

```


```

Class      1.000000
V11        0.154876
V4         0.133447


```

V2	0.091289
V21	0.040413
V19	0.034783
V20	0.020090
V8	0.019875
V27	0.017580
V28	0.009536
Amount	0.005632
V26	0.004455
V25	0.003308
V22	0.000805
V23	-0.002685
V15	-0.004223
V13	-0.004570
V24	-0.007221
Time	-0.012323
V6	-0.043643
V5	-0.094974
V9	-0.097733
V1	-0.101347
V18	-0.111485
V7	-0.187257
V3	-0.192961
V16	-0.196539
V10	-0.216883
V12	-0.260593
V14	-0.302544
V17	-0.326481

Name: Class, dtype: float64



```
[7]: # Heatmap with Class Correlation values only
plt.figure
sns.barplot(x = corr_with_class.index, y = corr_with_class.values, color = 'magenta')
plt.xticks(rotation = 90)
plt.title("Feature Correlation with Class (Fraud)")
plt.tight_layout()
plt.show()
```




### 2.0.1 Boxplot with Features that are most relevant to Fraud detection [Gives a good idea of what variables to include in our model generation]


```
[8]: # Compute correlations with Class
correlations = df.corr()['Class'].drop('Class').abs().sort_values(ascending =
↪ False)

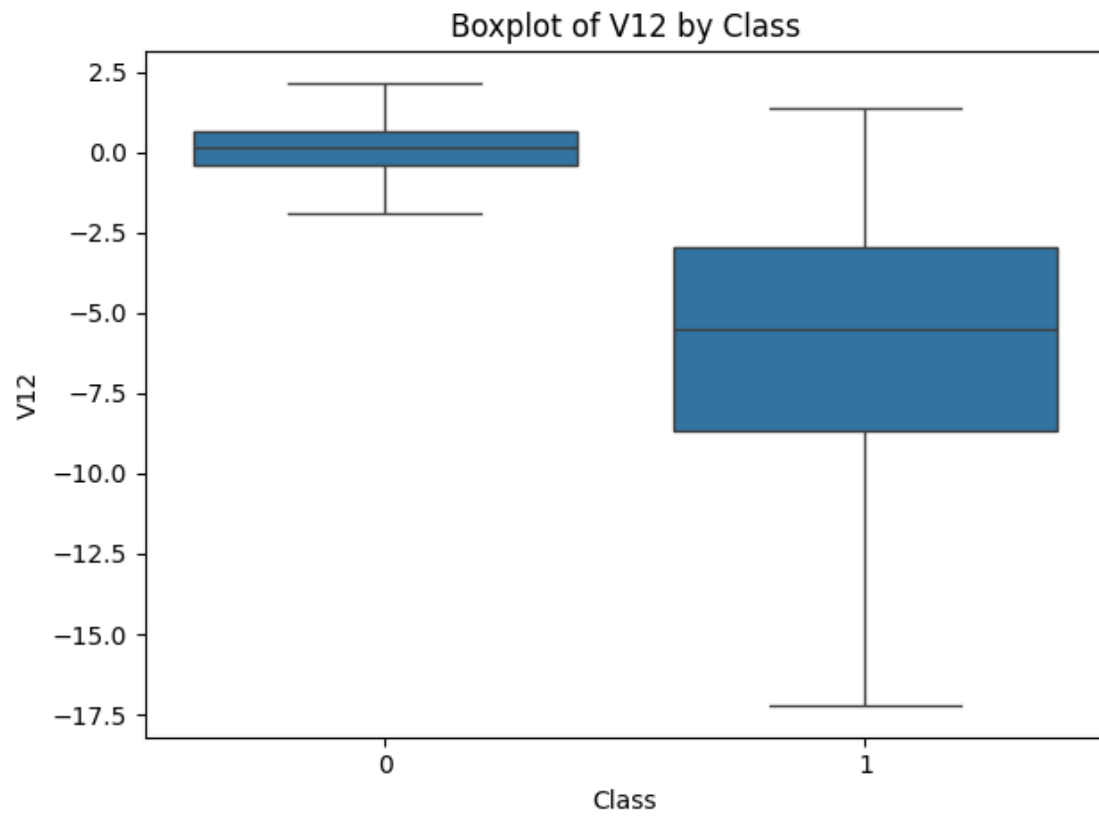
# Top 10 correlation
top_correlations = correlations.head(10).index


# Plot boxplot for top features
for column in top_correlations:
    sns.boxplot(x = 'Class', y = column, data = df, showfliers = False)
    plt.title(f'Boxplot of {column} by Class')
    plt.tight_layout()
    plt.show()
```
































```
[9]: sns.histplot(data = df, x = "Time", hue = "Class", bins = 50, element = "step",  
    stat = "density", common_norm = False)  
plt.title("Transaction Time Distribution by Class")  
plt.show()
```



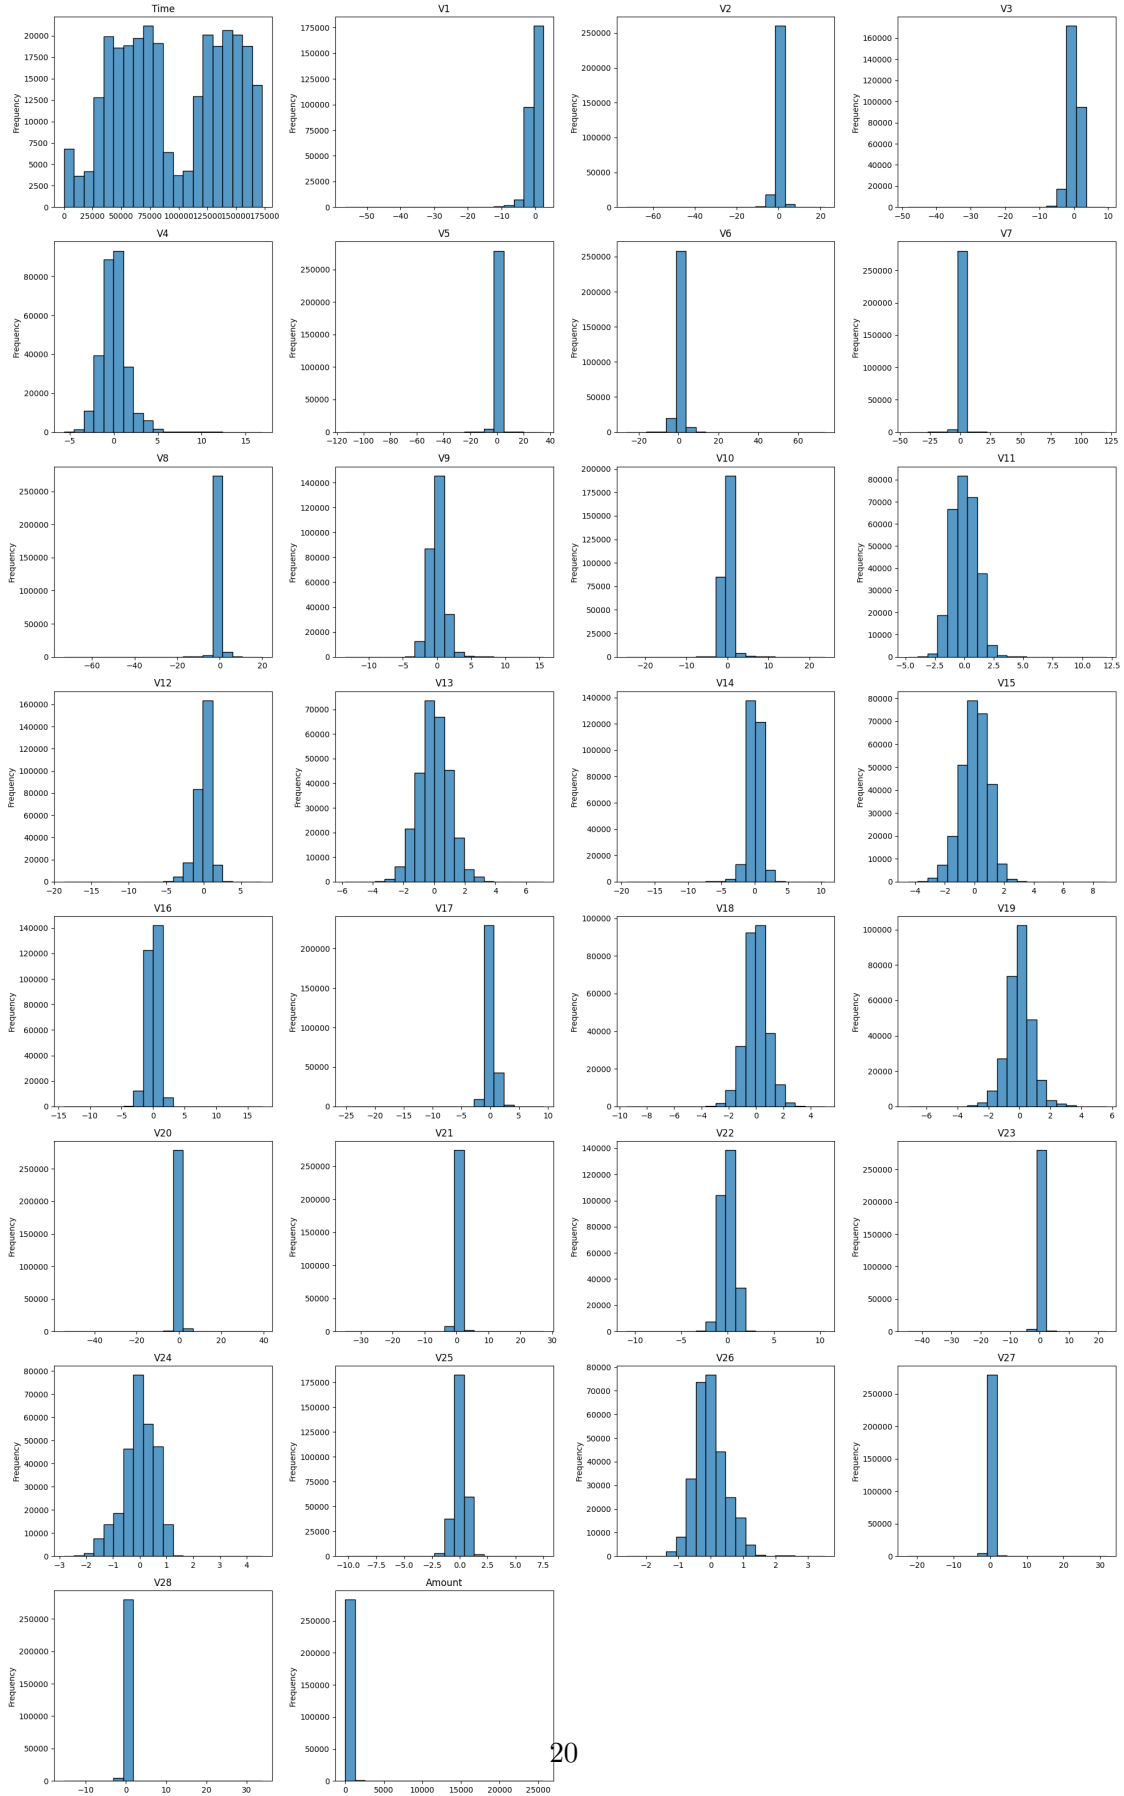
```
[10]: # Histogram to visualize centrality and distribution
columns_to_plot = df.drop('Class', axis = 1).columns
n_cols = 4
n_rows = int(np.ceil(len(columns_to_plot) / n_cols))

plt.figure(figsize = (n_cols * 5, n_rows * 4))

for i, var in enumerate(columns_to_plot):
    plt.subplot(n_rows, n_cols, i + 1)
    sns.histplot(df[var], kde = False, bins = 20)
    plt.title(f'{var}')
    plt.xlabel('')
    plt.ylabel('Frequency')

plt.tight_layout()
plt.suptitle('Histograms of All Features', fontsize = 18, y = 1.02)
plt.show()
```

# Histograms of All Features



## 2.1 SMOTE for Imbalanced Data

```
[11]: # Applying SMOTE for imbalanced class types
X = df.drop('Class', axis = 1)
y = df['Class']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
    random_state = 42, stratify = y)

# Apply smote
smote = SMOTE(random_state = 42)

X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

print(f'Original dataset shape: {y_train.value_counts()}')
print(f'Resampled dataset shape: {y_train_smote.value_counts()}')
```

```
Original dataset shape: Class
0    227451
1      394
Name: count, dtype: int64
Resampled dataset shape: Class
0    227451
1    227451
Name: count, dtype: int64
```

```
[12]: # Standardize features (important for SVM and MLP)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## 3 Model Strategies

### 3.1 Linear SVM

```
[13]: svm_model = LinearSVC(max_iter = 10000, random_state = 42)
svm_model.fit(X_train_scaled, y_train)

preds_svm = svm_model.predict(X_test_scaled)

print("Confusion Matrix:\n", confusion_matrix(y_test, preds_svm))
print(classification_report(y_test, preds_svm))
```

```
Confusion Matrix:
[[56852   12]
```

[ 40 58]]					
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	56864	
1	0.83	0.59	0.69	98	
accuracy			1.00	56962	
macro avg	0.91	0.80	0.85	56962	
weighted avg	1.00	1.00	1.00	56962	

```
[14]: # Linear SVM hyperparameter tuning
svm_tune = LinearSVC(max_iter = 10000, random_state = 42)

svm_params = {
    'C': uniform(0.001, 10),
    'class_weight': [None, 'balanced']
}

svm_search = RandomizedSearchCV(
    svm_tune,
    param_distributions = svm_params,
    n_iter = 20,
    scoring = 'f1',
    cv = 3,
    n_jobs = -1,
    verbose = 2,
    random_state = 42
)
svm_search.fit(X_train_scaled, y_train)

print("Best Linear SVM parameters:", svm_search.best_params_)
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits  
 Best Linear SVM parameters: {'C': np.float64(3.746401188473625), 'class\_weight': None}

```
[15]: # Retrain linear svm model
best_svm = LinearSVC(**svm_search.best_params_, random_state = 42, max_iter = 5000)
best_svm.fit(X_train_scaled, y_train)

preds_svm = best_svm.predict(X_test_scaled)

print("Confusion Matrix:\n", confusion_matrix(y_test, preds_svm))
print(classification_report(y_test, preds_svm))
```

Confusion Matrix:  
 [[56852 12]

```
[ 40 58]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    56864
     1       0.83      0.59      0.69      98

 accuracy          1.00          1.00    56962
 macro avg       0.91      0.80      0.85    56962
weighted avg       1.00      1.00      1.00    56962
```

### 3.2 Neural Network (MLP)

```
[16]: mlp_model = MLPClassifier(
        hidden_layer_sizes = (50,),
        max_iter = 500,
        random_state = 42,
        early_stopping = True
    )
mlp_model.fit(X_train_scaled, y_train)

probs_mlp = mlp_model.predict_proba(X_test_scaled)[: , 1]
preds_mlp = mlp_model.predict(X_test_scaled)
# preds_mlp = (probs_mlp >= 0.5).astype(int)

print("AUC:", roc_auc_score(y_test, probs_mlp))
print("Confusion Matrix:\n", confusion_matrix(y_test, preds_mlp))
print(classification_report(y_test, preds_mlp))
```

AUC: 0.9805382409013127

Confusion Matrix:

```
[[56855 9]
 [ 25 73]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00    56864
     1       0.89      0.74      0.81      98

 accuracy          1.00          1.00    56962
 macro avg       0.94      0.87      0.91    56962
weighted avg       1.00      1.00      1.00    56962
```

```
[17]: # Neural network hyperparameter tuning
mlp_tune = MLPClassifier(max_iter = 300, random_state = 42, early_stopping = True)

mlp_params = {
```

```

    'hidden_layer_sizes': [(50,), (100,), (50, 25)],
    'activation': ['relu', 'tanh'],
    'alpha': [0.0001, 0.001, 0.01],
    'learning_rate_init': [0.001, 0.01]
}

mlp_search = RandomizedSearchCV(
    mlp_tune,
    param_distributions = mlp_params,
    scoring = 'f1',
    cv = 3,
    n_jobs = -1,
    verbose = 2
)
mlp_search.fit(X_train_scaled, y_train)

print("Best MLP parameters:", mlp_search.best_params_)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

Best MLP parameters: {'learning\_rate\_init': 0.001, 'hidden\_layer\_sizes': (50,), 'alpha': 0.001, 'activation': 'tanh'}

```

[18]: # Retrain neural net model
best_mlp = MLPClassifier(**mlp_search.best_params_, random_state = 42,
    ↪early_stopping = True)
best_mlp.fit(X_train_scaled, y_train)

probs_mlp = best_mlp.predict_proba(X_test_scaled)[: , 1]
preds_mlp = best_mlp.predict(X_test_scaled)

print("AUC:", roc_auc_score(y_test, probs_mlp))
print("Confusion Matrix:\n", confusion_matrix(y_test, preds_mlp))
print(classification_report(y_test, preds_mlp))

```

AUC: 0.9768028694313966

Confusion Matrix:

```

[[56851   13]
 [   18   80]]

```

		precision	recall	f1-score	support
	0	1.00	1.00	1.00	56864
	1	0.86	0.82	0.84	98
accuracy				1.00	56962
macro avg		0.93	0.91	0.92	56962
weighted avg		1.00	1.00	1.00	56962



### 3.3 Naive Bayes (Gaussian)

```
[19]: nb_model = GaussianNB()
nb_model.fit(X_train_scaled, y_train)

probs_nb = nb_model.predict_proba(X_test_scaled)[: , 1]
preds_nb = nb_model.predict(X_test_scaled)

print("AUC:", roc_auc_score(y_test, probs_nb))
print("Confusion Matrix:\n", confusion_matrix(y_test, preds_nb))
print(classification_report(y_test, preds_nb))
```

AUC: 0.963247971529636

Confusion Matrix:

```
[[55535  1329]
```

```
[   15    83]]
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	56864
1	0.06	0.85	0.11	98
accuracy			0.98	56962
macro avg	0.53	0.91	0.55	56962
weighted avg	1.00	0.98	0.99	56962

```
[20]: # Naive bayes hyperparameter tuning
nb_params = {
    'var_smoothing': np.logspace(-9, -2, 8)
}

nb_search = RandomizedSearchCV(
    GaussianNB(),
    param_distributions = nb_params,
    n_iter = 8,
    scoring = 'f1',
    cv = 3,
    n_jobs = -1,
    verbose = 2,
    random_state = 42
)
nb_search.fit(X_train_scaled, y_train)

print("Best NB parameters:", nb_search.best_params_)
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

Best NB parameters: {'var\_smoothing': np.float64(0.01)}

```
[21]: # Retrain naive bayes model
best_nb = GaussianNB(**nb_search.best_params_)
best_nb.fit(X_train_scaled, y_train)

probs_nb = best_nb.predict_proba(X_test_scaled)[: , 1]
preds_nb = best_nb.predict(X_test_scaled)

print("AUC:", roc_auc_score(y_test, probs_nb))
print("Confusion Matrix:\n", confusion_matrix(y_test, preds_nb))
print(classification_report(y_test, preds_nb))
```

AUC: 0.9632567644390339

Confusion Matrix:

```
[[55548  1316]
 [   15    83]]
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	56864
1	0.06	0.85	0.11	98
accuracy			0.98	56962
macro avg	0.53	0.91	0.55	56962
weighted avg	1.00	0.98	0.99	56962

### 3.4 Logistic regression (baseline)

```
[22]: def best_threshold_by_f1(y_true, scores):
    precision, recall, thresholds = precision_recall_curve(y_true, scores)
    thresholds = np.append(thresholds, 1.0)
    f1 = 2 * (precision * recall) / (precision + recall + 1e-12)
    i = int(np.nanargmax(f1))
    return float(thresholds[i]), float(f1[i]), float(precision[i]),
    ↪float(recall[i])

def evaluate(name, y_true, scores):
    roc = roc_auc_score(y_true, scores)
    pr = average_precision_score(y_true, scores)
    thr, f1, prec, rec = best_threshold_by_f1(y_true, scores)
    return {
        "Model": name,
        "ROC_AUC": roc,
        "PR_AUC": pr,
        "Best_F1": f1,
        "Best_Thresh": thr,
        "Precision@BestF1": prec,
        "Recall@BestF1": rec,
```

```

    }

results = []

# Logistic Regression
logreg_model = Pipeline([
    ("scaler", StandardScaler()),
    ("clf", LogisticRegression(max_iter = 2000, class_weight = "balanced",
    ↪ solver = "lbfgs")),
])
logreg_model.fit(X_train, y_train)
logreg_scores = logreg_model.predict_proba(X_test)[: , 1]
results.append(evaluate("Logistic Regression (baseline)", y_test,
    ↪ logreg_scores))

```

### 3.5 Decision Tree

```

[23]: # Decision Tree
dt_model = DecisionTreeClassifier(
    random_state = 42,
    class_weight = "balanced",
    min_samples_leaf = 2)
dt_model.fit(X_train, y_train)
dt_scores = dt_model.predict_proba(X_test)[: , 1]
results.append(evaluate("Decision Tree", y_test, dt_scores))

```

### 3.6 Random Forest

```

[24]: # Random Forest
rf_model = RandomForestClassifier(
    n_estimators = 400,
    random_state = 42,
    n_jobs = -1,
    class_weight = "balanced_subsample",
    min_samples_leaf = 2
)
rf_model.fit(X_train, y_train)
rf_scores = rf_model.predict_proba(X_test)[: , 1]
results.append(evaluate("Random Forest", y_test, rf_scores))

```

### 3.7 XGBoost/LightGBM/fallback

```

[25]: # XGBoost / LightGBM / fallback
pos, neg = (y_train==1).sum(), (y_train==0).sum()
scale_pos_weight = neg / max(pos, 1)

try:

```

```

xgb = XGBClassifier(
    random_state = 42,
    n_estimators = 500,
    max_depth = 3,
    learning_rate = 0.05,
    subsample = 0.8,
    colsample_bytree = 0.8,
    reg_lambda = 1.0,
    scale_pos_weight = scale_pos_weight,
    tree_method = "hist",
    eval_metric = "logloss"
)
xgb.fit(X_train, y_train)
xgb_scores = xgb.predict_proba(X_test)[:, 1]
results.append(evaluate("XGBoost", y_test, xgb_scores))
except Exception:
    try:
        lgbm = LGBMClassifier(
            random_state = 42,
            n_estimators = 800,
            learning_rate = 0.05,
            num_leaves = 31,
            subsample = 0.8,
            colsample_bytree = 0.8,
            reg_lambda = 1.0,
            scale_pos_weight = scale_pos_weight
        )
        lgbm.fit(X_train, y_train)
        lgbm_scores = lgbm.predict_proba(X_test)[:, 1]
        results.append(evaluate("LightGBM", y_test, lgbm_scores))
    except Exception:
        print("XGBoost/LightGBM not available. Using HistGradientBoosting as
↳fallback.")
        hgb = HistGradientBoostingClassifier(max_iter = 500, learning_rate = 0.
↳05, random_state = 42)
        hgb.fit(X_train, y_train)
        try:
            hgb_scores = hgb.predict_proba(X_test)[:, 1]
        except Exception:
            raw = hgb.decision_function(X_test)
            hgb_scores = 1 / (1 + np.exp(-raw))
        results.append(evaluate("HistGradientBoosting (fallback)", y_test,
↳hgb_scores))

```

[26]: *# Evaluate tuned models*

```

svm_scores = best_svm.decision_function(X_test_scaled)
results.append(evaluate("Linear SVM", y_test, svm_scores))

```

```

mlp_scores = best_mlp.predict_proba(X_test_scaled)[: , 1]
results.append(evaluate("Neural Network (MLP)", y_test, mlp_scores))

nb_scores = best_nb.predict_proba(X_test_scaled)[: , 1]
results.append(evaluate("Naive Bayes (Gaussian)", y_test, nb_scores))

# Results
results_df = pd.DataFrame(results).sort_values("PR_AUC", ascending = False).
    ↪reset_index(drop = True)
print(results_df.to_string(index = False))

```

	Model	ROC_AUC	PR_AUC	Best_F1	Best_Thresh
Precision@BestF1	Recall@BestF1				
	Random Forest	0.964768	0.854731	0.874317	0.412421
0.941176	0.816327				
	Neural Network (MLP)	0.976803	0.853069	0.842105	0.535248
0.869565	0.816327				
	XGBoost	0.984078	0.846238	0.813953	0.998153
0.945946	0.714286				
	Linear SVM	0.942820	0.749218	0.818653	-0.325878
0.831579	0.806122				
	Logistic Regression (baseline)	0.972083	0.718971	0.824742	1.000000
0.833333	0.816327				
	Decision Tree	0.861951	0.497390	0.692683	0.998271
0.663551	0.724490				
	Naive Bayes (Gaussian)	0.963257	0.082185	0.168264	1.000000
0.093936	0.806122				

### 3.8 ROC curves of all models

```


[27]: plt.figure(figsize = (8,6))

for name, scores in [
    ("Random Forest", rf_scores),
    ("Neural Network (MLP)", mlp_scores),
    ("XGBoost", xgb_scores),
    ("Linear SVM", svm_scores),
    ("Naive Bayes", nb_scores)
]:
    fpr, tpr, _ = roc_curve(y_test, scores)
    auc = roc_auc_score(y_test, scores)
    plt.plot(fpr, tpr, label = f"{name} (AUC = {auc:.3f})")

plt.plot([0,1], [0,1], 'k--', label = "Random")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")

```

```
plt.title("ROC Curves for All Models")
plt.legend()
plt.grid(True)
plt.show()
```



### 3.9 Precision-Recall Curves (Top 3 Models)

```
[28]: plt.figure(figsize = (8,6))

for name, scores in [
    ("Random Forest", rf_scores),
    ("Neural Network (MLP)", mlp_scores),
    ("XGBoost", xgb_scores)
]:
    precision, recall, _ = precision_recall_curve(y_test, scores)
    plt.plot(recall, precision, label = name)

plt.xlabel("Recall")
plt.ylabel("Precision")
```

```
plt.title("Precision-Recall Curves for Top 3 Models")
plt.legend()
plt.grid(True)
plt.show()
```

